



Linear-Time Sorting Algorithms

Sorting So Far

- Insertion sort:
 - Easy to code
 - Fast on small inputs (less than ~50 elements)
 - Fast on nearly-sorted inputs
 - $O(n^2)$ worst case
 - $O(n^2)$ average (equally-likely inputs) case
 - $O(n^2)$ reverse-sorted case

Sorting So Far

- Merge sort:
 - Divide-and-conquer:
 - ◆ Split array in half
 - ◆ Recursively sort subarrays
 - ◆ Linear-time merge step
 - $O(n \lg n)$ worst case
 - Doesn't sort in place

Sorting So Far

- Heap sort:
 - Uses the very useful heap data structure
 - ◆ Complete binary tree
 - ◆ Heap property: parent key $>$ children's keys
 - $O(n \lg n)$ worst case
 - Sorts in place

Sorting So Far

- Quick sort:
 - Divide-and-conquer:
 - ◆ Partition array into two subarrays, recursively sort
 - ◆ All of first subarray $<$ all of second subarray
 - ◆ No merge step needed!
 - $O(n \lg n)$ average case
 - Fast in practice
 - $O(n^2)$ worst case
 - ◆ Naïve implementation: worst case on sorted input

Comparison sort

- Comparison sort:
 - Insertion sort, $O(n^2)$, upper bound in worst case
 - Merge sort, $O(n \lg n)$, upper bound in worst case
 - Heapsort, $O(n \lg n)$, upper bound in worst case
 - Quicksort, $O(n \lg n)$, in average case
- Question:
 - what is the lower bounds for any comparison sorting: i.e., at least how many comparisons needed in worst case?
 - It turns out: Lower bound in worst case: $\Omega(n \lg n)$, how to prove?
 - Merge and Heapsort are asymptotically optimal comparison sorts.

How Fast Can We Sort?

- We will provide a lower bound, then beat it
 - *How do you suppose we'll beat it?*
- First, an observation: all of the sorting algorithms so far are *comparison sorts*
 - The only operation used to gain ordering information about a sequence is the pairwise comparison of two elements

Sorting In Linear Time

- Counting sort
 - No comparisons between elements!
 - *But*...depends on assumption about the numbers being sorted
 - ◆ We assume numbers are in the range $1..k$
 - The algorithm:
 - ◆ Input: $A[1..n]$, where $A[j] \in \{1, 2, 3, \dots, k\}$
 - ◆ Output: $B[1..n]$, sorted (notice: not sorting in place)
 - ◆ Also: Array $C[1..k]$ for auxiliary storage

Counting Sort

```
1      CountingSort(A, B, k)
2          for i=1 to k
3              C[i]= 0;
4          for j=1 to n
5              C[A[j]] += 1;
6          for i=2 to k
7              C[i] = C[i] + C[i-1];
8          for j=n downto 1
9              B[C[A[j]]] = A[j];
10         C[A[j]] -= 1;
```

Work through example: $A=\{4\ 1\ 3\ 4\ 3\}$, $k = 4$

Counting Sort

```
1  CountingSort(A, B, k)
2      for i=1 to k
3          C[i]= 0;
4      for j=1 to n
5          C[A[j]] += 1;
6      for i=2 to k
7          C[i] = C[i] + C[i-1];
8      for j=n downto 1
9          B[C[A[j]]] = A[j];
10         C[A[j]] -= 1;
```

Takes time $O(k)$

Takes time $O(n)$

What will be the running time?

Counting Sort

- Total time: $O(n + k)$
 - Usually, $k = O(n)$
 - Thus counting sort runs in $O(n)$ time
- But sorting is $\Omega(n \lg n)$!
 - No contradiction--this is not a comparison sort (in fact, there are *no* comparisons at all!)
 - Notice that this algorithm is *stable*

Counting Sort

- Cool! *Why don't we always use counting sort?*
- Because it depends on range k of elements
- *Could we use counting sort to sort 32 bit integers? Why or why not?*
- Answer: no, k too large ($2^{32} = 4,294,967,296$)